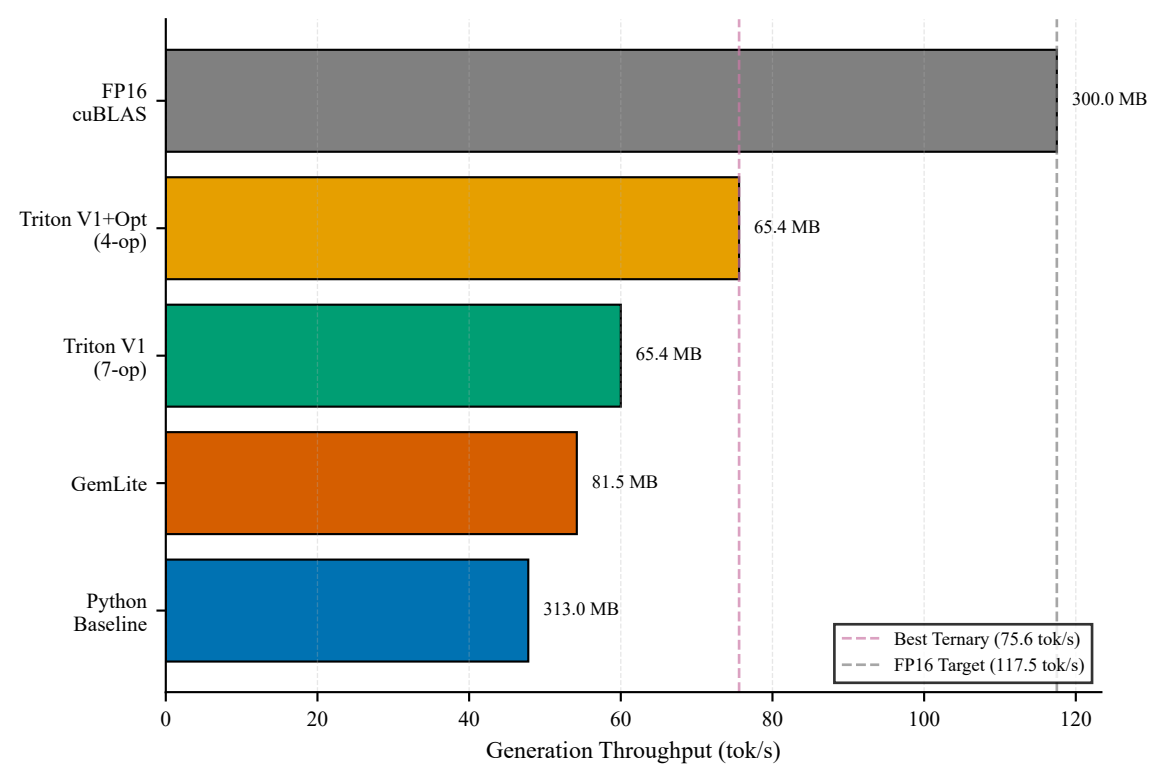
FP16
cuBLAS
300.0 MB
Triton V1+Opt
(4-op)
65.4 MB
Triton V1
(7-op)
65.4 MB
GemLite
81.5 MB
Python
Baseline
313.0 MB
Best Ternary (75.6 tok/s)
FP16 Target (117.5 tok/s)
0
20
40
60
80
100
120
Generation Throughput (tok/s)