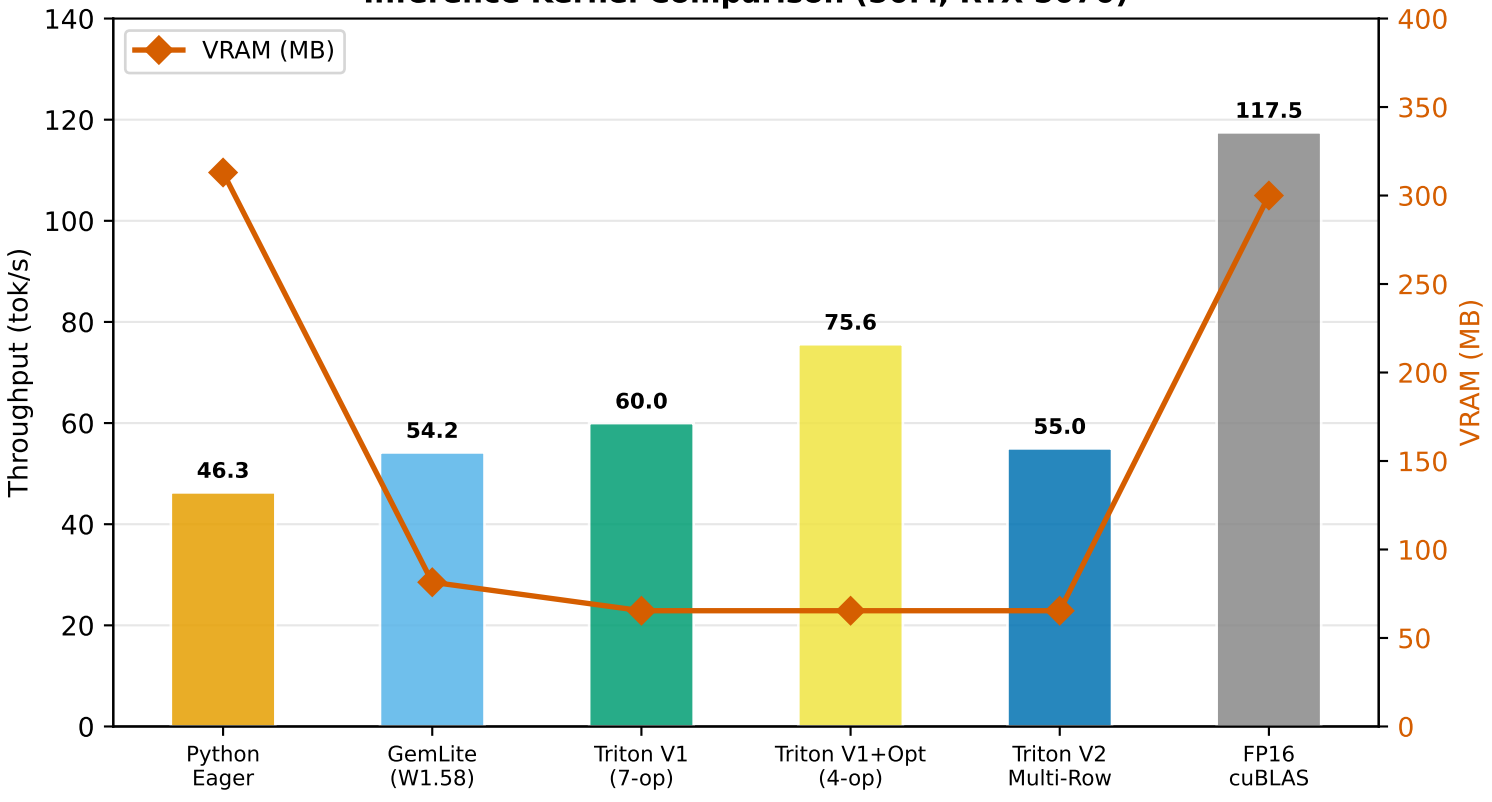
Inference Kernel Comparison (50M, RTX 5070)
VRAM (MB)
Throughput (tok/s)
VRAM (MB)
0
20
40
60
80
100
120
140
0
50
100
150
200
250
300
350
400
46.3
54.2
60.0
75.6
55.0
117.5
Python Eager
GemLite (W1.58)
Triton V1 (7-op)
Triton V1+Opt (4-op)
Triton V2 Multi-Row
FP16 cuBLAS